

Decipher Delta Lake

By Nisha Kumari and Divya Dua


ThoughtWorks®

© 2020 ThoughtWorks

Agenda

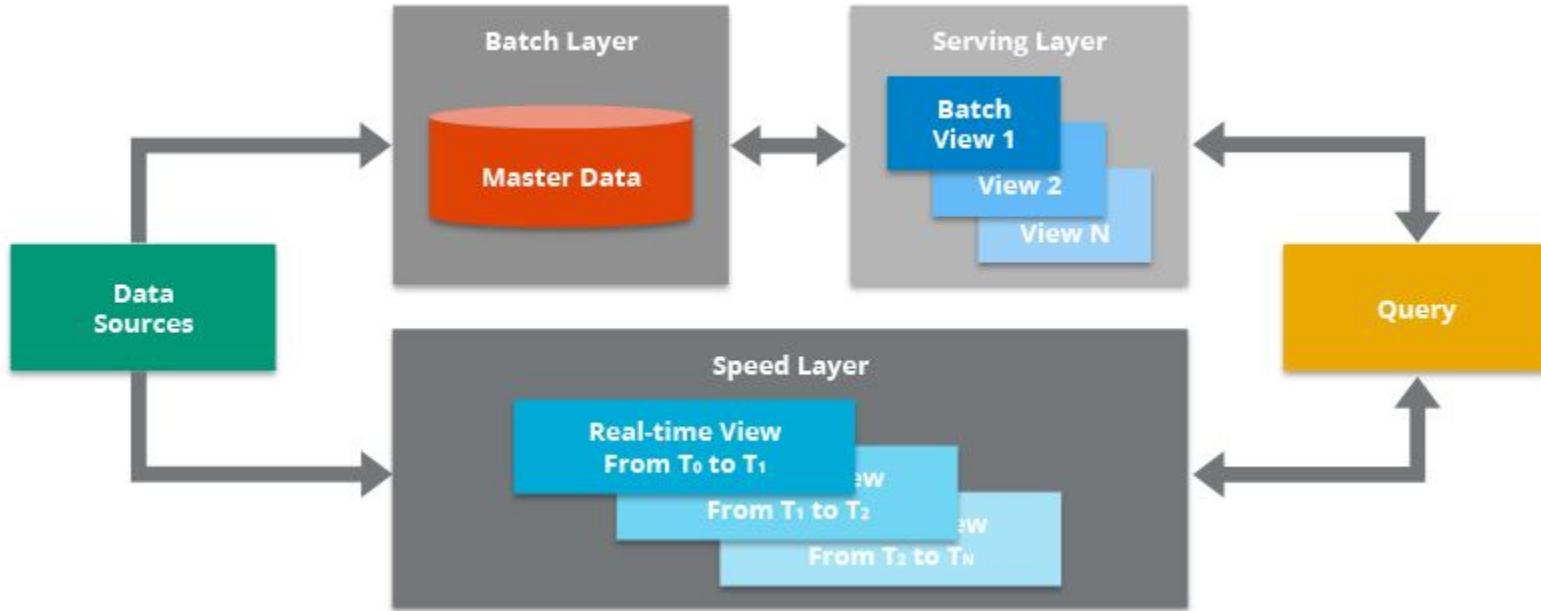
- **Towards Delta**
 - Traditional Data Lakes
 - Lambda architecture
 - Current challenges
- **Delta Lake**
 - Why delta?
 - Architecture
 - Delta tables
- **Features of Delta Engines**
 - ACID Transactions
 - Schema enforcement & evolution
 - Time Travel
 - Performance Optimization, Zorder etc

Traditional Data Lake (Handling Volume)



1. Input - Data from many sources.
2. Collected in a data lake.
3. Processed using some distributed processing engine.
4. Served for analysis.

Lambda Architecture (Handling Velocity)



Challenges with current architecture

1. Query Performance

- ETL process can add latency
- With growing scale, query becomes slower.

2. System Complexity

- Unifying stream and batch doubles infra cost.
- Low level Code intervention needed to sync two pipeline data.
- Need to reset pipelines for stream failures.

3. Data Reliability

- Failed jobs can corrupt data
- Schema changes can break joins, aggregates, transformations.
- Hard to update/ delete data
- Concurrent access suffer inconsistent query results.

Delta to the rescue!

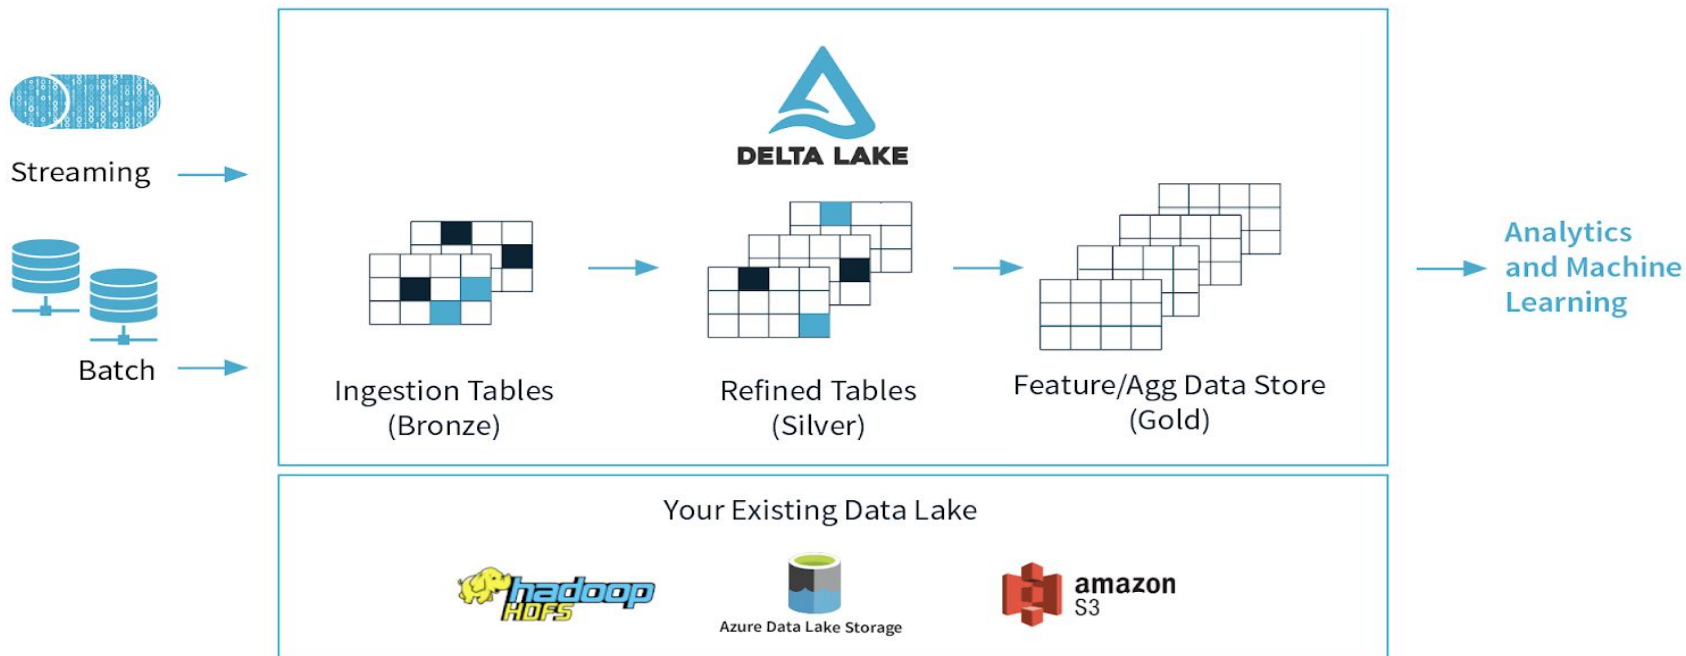


<https://blog.knoldus.com/why-databricks-delta-called-as-unified-data-management/>

- Delta stores data in versioned **Apache Parquet format**.
- Uses **Delta Transaction log protocol** to provide consistency.
- Provide Automatic **data indexing**.

*Incoming data is processed as “**delta**” records rather than the append-only new records.*


Delta Lake Architecture



Delta Table

To provide ACID guarantees we require the storage system to provide the following.

1. **Atomic visibility:** There must be a way for a file to be visible in its entirety or not visible at all.
2. **Mutual exclusion:** Only one writer must be able to create (or rename) a file at the final destination.
3. **Consistent listing:** Once a file has been written in a directory, all future listings for that directory must return that file.



<https://github.com/delta-io/delta/blob/master/PROTOCOL.md>

Storage systems do not necessarily provide all of these guarantees out-of-the-box.

Delta Lake transactional operations typically go through the **LogStore API** instead of accessing the storage system directly.

Delta Lake Features

Delta Lake is an open source storage layer that brings reliability to data lakes.

1. **ACID transactions**
2. **Scalable metadata handling**
3. **Time travel/ Data versioning**
4. **Streaming and batch unification**
5. **Schema enforcement**
6. **Schema evolution**
7. **Updates and deletes**
8. **100% compatible with Apache spark**

Dive into Delta Lake!

1. **ACID Transactions Demo**
2. **Schema Enforcement and Evolution Demo**

Demo Summary

1. Delta table is a **single serial history of atomic versions**.
2. The state of a table at a given version is called a **snapshot**.
3. Delta Lake uses ***optimistic concurrency control*** for concurrent read-write access.

```
/mytable/_delta_log/00000000000000000000.json  
/mytable/_delta_log/00000000000000000001.json  
/mytable/_delta_log/00000000000000000003.json  
/mytable/_delta_log/00000000000000000003.checkpoint.parquet  
/mytable/_delta_log/_last_checkpoint  
/mytable/part-00000-3935a07c-416b-4344-ad97-2a38342ee2fc.c000.snappy.parquet
```

Data reliability by Delta

- **ACID Transactions** : “all or nothing” transactions
- **Snapshot Isolation** : multiple readers & writers support
- **Schema Enforcement**: Delta provides schema and prevent writes that do not align
- **Schema Evolution**: via merge schema or overwrite schema as needed.
- **Exactly Once**: Delta employs checkpointing to ensure data is not missed or repeated.
- **Upserts and Deletes support**: Spark tables are write once, modifications need to be done explicitly, while delta has implicit support for them.

Data Versioning

We have seen Delta Transactions & its Logs. Now, lets see how to use them for data versioning?

- Query older snapshot of delta table
- Revert to earlier versions of data for audits
- Trace transactions on the delta lake
- Reproducible machine learning experiments



<https://www.thesun.co.uk/tech/8624077/time-machine1experiment-quantum-physics-electrons/>

Performance Optimization

So far, delta lake supported us with ACID, data versioning, schema enforcement, etc.

But,

Is there any downside of this architecture?

- 1. Query times?**

- a. Delta scans so many files

- 2. Storage?**

- a. Delta maintain all versions of data

Now, lets see how delta helps improve query times.

Optimize via File Management

Data Compaction or Bin-Packing


- Coalesce small files into larger ones.
- Produce evenly-balanced data files w.r.t their size on disk.
- No data related changes to the table.
- Idempotent.

Data Skipping

- Automatically collect statistics when you write data into a Delta table.
 - Parquet metadata supports this.
- Improve query execution time by skipping unrelated files.

For best results, apply **Z-Ordering**. ()

Z-ordering (multi-dimensional clustering)



- **Dot** = data points
- **Box** = data files; in this example, we aim for files of 4 points each
- **Yellow boxes** = data file that's read for the given query
- **Green dot** = data point that passes the query's filter and answers the query
- **Red dot** = data point that's read, but doesn't satisfy the filter; "false positive"

vacuum

- Remove files older than the retention threshold. Default is 7 days.
- Deletes only ***data files, not log files.***
- Not triggered automatically.
- Ability to time travel back to a version older than retention period is lost after running vacuum.

Delta Cache

- Accelerates data reads by creating copies of remote files in nodes' local storage.
- Supports only Parquet files, not other storage formats.
- Stores data entirely on the local disk.
- Any stale entries are automatically invalidated and evicted from the cache.
- uses efficient ***decompression algorithms*** and outputs data in the optimal format

References

- <https://docs.delta.io/latest/delta-batch.html#language-scala> - Demo
- <https://docs.microsoft.com/en-us/azure/databricks/delta/>
- <https://akashrehan.wordpress.com/2019/07/11/anatomy-of-databricks-delta-lake/>
- <https://blog.knoldus.com/databricks-delta-architecture/>
- <https://github.com/delta-io/delta>

Thank You

Divya Dua

Data Engineer

divya.dua@thoughtworks.com

Nisha Kumari

Data Engineer

nishak@thoughtworks.com

ThoughtWorks®