

Decipher Delta Lake

By Nisha Kumari and Divya Dua

ThoughtWorks®

© 2020 ThoughtWorks

Agenda

- **Intro**


- **Data lake**
- **Challenges & Mitigation**
- **Possible solution**
- **Delta Lake**
- **Delta Lake architecture**

- **Demo**

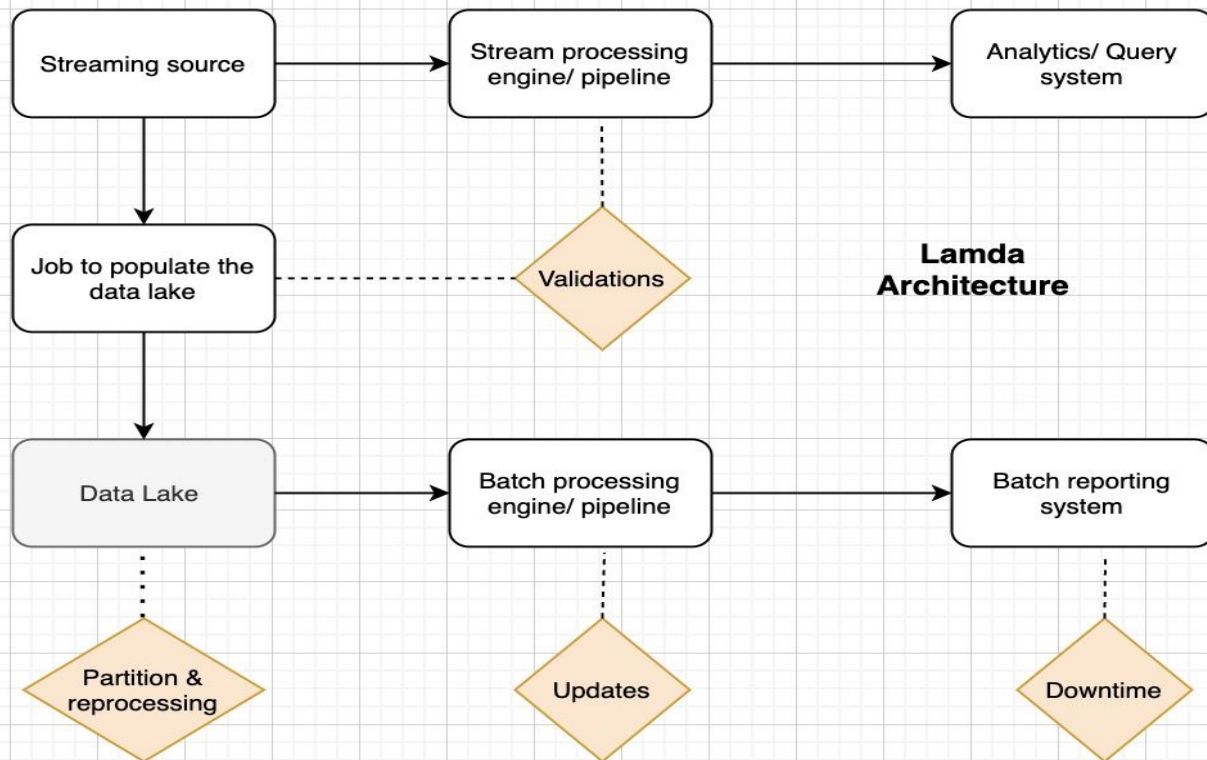
- **Setup**
- **ACID**
- **Schema evolution**
- **Data versioning**

- **Performance tuning**

Data Lake



Challenges & Mitigations



1. Data velocity
2. Historical queries
3. Messy Data
4. Job failures
5. Updates

Possible solution

A storage solution that can provide transactional guarantees.

1. **Atomic visibility:** There must be a way for a file to be visible in its entirety or not visible at all.
2. **Mutual exclusion:** Only one writer must be able to create (or rename) a file at the final destination.
3. **Consistent listing:** Once a file has been written in a directory, all future listings for that directory must return that file.


Typical Storage systems do not necessarily provide this. This is where the **delta lake** comes in play.

*Delta Lake transactional operations typically go through the **LogStore API** instead of accessing the storage system directly. (called Delta Transaction Log protocol).*


Delta Lake

1. Storage Layer
2. Open source
3. Based on Spark APIs
 - a. Stream and batch
 - b. First class support for upserts and deletes.
4. Uses parquet as underlying filesystem.
5. Uses Delta Transaction Log Protocol. Logs
 - a. committed before any modification.
 - b. maintain data change.
 - c. maintain & enforce schema.
 - d. Allow time travel.
6. Scalable metadata handling & checkpointing.

Delta Lake

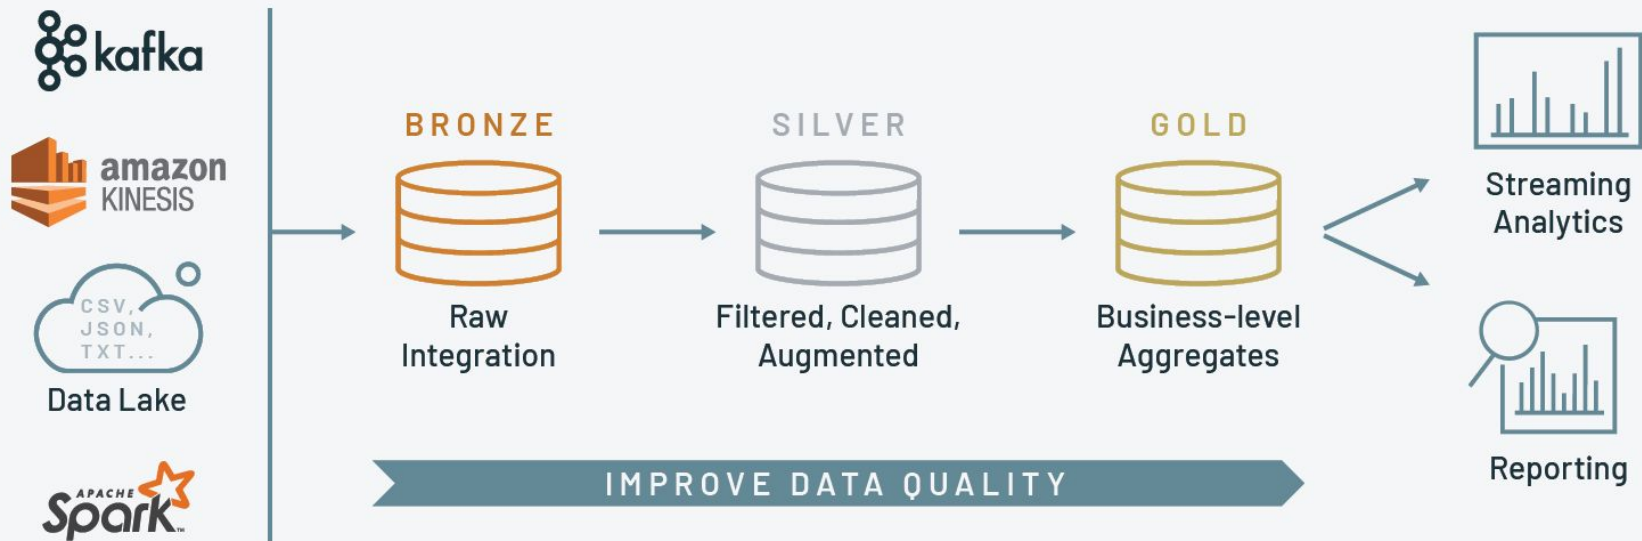


parquet files



transaction logs

Delta Lake Architecture



<https://databricks.com/blog/2020/06/02/monitor-your-databricks-workspace-with-audit-logs.html>

Demo

1. Setup
2. ACID Transactions
3. Data Versioning
4. Schema evolution

Performance tuning

1. On schema aspect
 - a. Uses Spark
 - b. Checkpointing
2. On data aspect
 - a. Compaction (# optimize)
 - b. Delete old files (# vacuum)
 - c. Data skipping (runs when applicable)
 - d. Colocating predicates (Zorder)

References

- <https://docs.delta.io/latest/delta-batch.html#language-scala> - Demo
- <https://docs.microsoft.com/en-us/azure/databricks/delta/>
- <https://akashrehan.wordpress.com/2019/07/11/anatomy-of-databricks-delta-lake/>
- <https://blog.knoldus.com/databricks-delta-architecture/>
- <https://github.com/delta-io/delta>

Thank You

Divya Dua

Data Engineer

divya.dua@thoughtworks.com

Nisha Kumari

Data Engineer

nishak@thoughtworks.com

ThoughtWorks®